

Boosting algorithm

Diane Lingrand and Frédéric Precioso



2020 - 2021

- Classification or regression algorithms
 - SVM, ANN, Decision Tree, Bayes Networks
- Ensemble algorithms
 - voting or averaging
 - bootstrapping - features sampling - Random Forests
 - boosting : today's topic
 - stacking

1 Boosting algorithm

- Idea
- A first simple example to understand boosting idea
- Probabilistic boosting and Adaboost
- A simple example for understanding Adaboost

2 Application to face detection

1 Boosting algorithm

- Idea
- A first simple example to understand boosting idea
- Probabilistic boosting and Adaboost
- A simple example for understanding Adaboost

2 Application to face detection

1 Boosting algorithm

■ Idea

- A first simple example to understand boosting idea
- Probabilistic boosting and Adaboost
- A simple example for understanding Adaboost

2 Application to face detection

- Supervised learning
 - learning dataset
 - test dataset
- Binary classification

Binary classification

- Separation between data corresponding to some criterions and data not corresponding.



Main idea of boosting

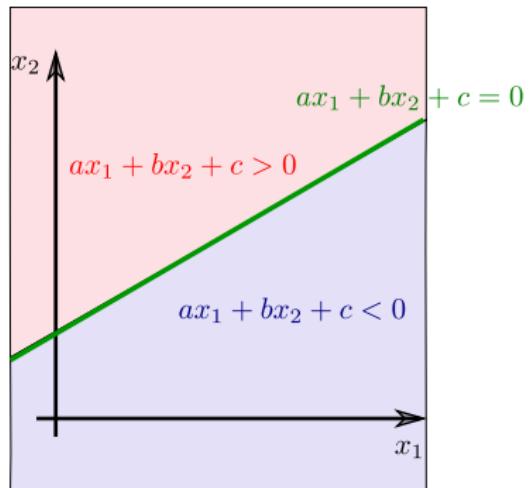
- Build a strong learner from weak learners
- Example :
 - Green corresponds to fir tree.
 - Vertical shapes corresponds to fir tree.
 - Shapes with bottom larger than top corresponds to fir tree.
 - When dominant color is red, it is not a fir tree.
 -
- Example of sport bet and heuristics.
 - We ask experts simple rules that are true more than 50%
 - We focus on examples for which a rule fails and we ask other rules to experts.
 - And loop the process

Origin of boosting

- Mathematician Kearns asks : “Is it possible to make as good as possible a weak learner” (that is to say better than random) ?
- Schapire, answered in 90 : “ Yes ! ” , and exhibited the first elementary boosting algorithm which shows that a weak binary classifier can always improve by being trained on 3 subsamples well chosen.
- The choice of the weak learner does not have any importance (a decision tree, a bayesian classifier, a SVM, a Neural Network, etc.), but one has to choose the 3 training subsamples with respect to its performance.

An example for a weak classifier : line in 2d plane

- points lie in a 2d plane
- the weak classifier is defined by a line of equation $ax_1 + bx_2 + c = 0$
 - divides the plane into 2 area :
 - $(x_1; x_2)$ for which $ax_1 + bx_2 + c > 0$
 - $(x_1; x_2)$ for which $ax_1 + bx_2 + c < 0$
- learning means to find coefficients a , b , and c .



1 Boosting algorithm

- Idea
- A first simple example to understand boosting idea
- Probabilistic boosting and Adaboost
- A simple example for understanding Adaboost

2 Application to face detection

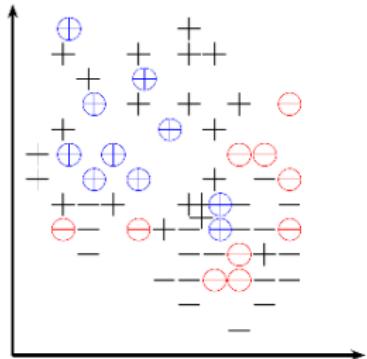
Elementary boosting algorithm

S : learning data set of m elements

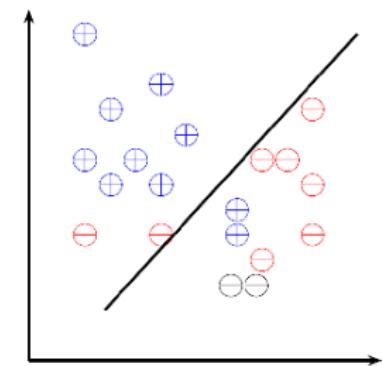
- ① Learn the classifier h_1 on subset $S_1 \subset S$. Test of h_1 on $S \setminus S_1$.
- ② Learn the classifier h_2 on subset $S_2 \subset S \setminus S_1$ with half of elements of S_2 wrongly classified by h_1 .
- ③ Learn classifier h_3 on subset $S_3 \subset S \setminus S_1 \setminus S_2$: contains elements for which rules h_1 and h_2 answer differently.
- ④ H : Majority vote between answers of h_1 , h_2 and h_3 .

A toy example (1)

Classifiers are lines. Learning a classifier corresponds to find the linear separation of data.



S : set of + and -
 S_1 : reds and blues



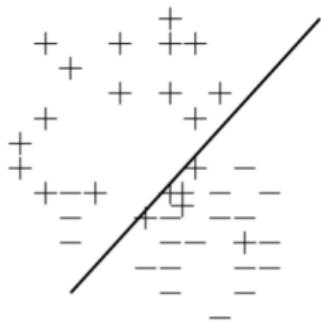
Classifier h_1 learned on S_1 .

Elementary boosting algorithm

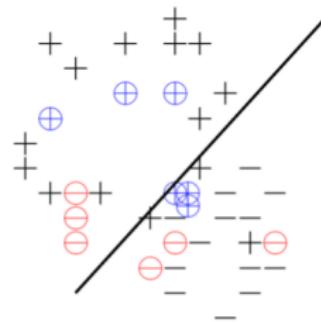
S : learning data set of m elements

- ① Learn the classifier h_1 on subset $S_1 \subset S$. Test of h_1 on $S \setminus S_1$.
- ② Learn the classifier h_2 on subset $S_2 \subset S \setminus S_1$ with half of elements of S_2 wrongly classified by h_1 .
- ③ Learn classifier h_3 on subset $S_3 \subset S \setminus S_1 \setminus S_2$: contains elements for which rules h_1 and h_2 answer differently.
- ④ H : Majority vote between answers of h_1 , h_2 and h_3 .

A toy example (2)



$S \setminus S_1$ and h_1



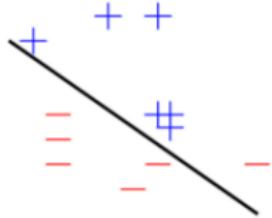
$S_2 \subset S \setminus S_1$: reds and blues

Elementary boosting algorithm

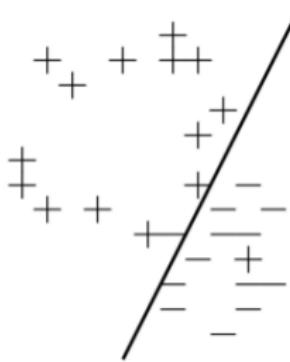
S : learning data set of m elements

- ① Learn the classifier h_1 on subset $S_1 \subset S$. Test of h_1 on $S \setminus S_1$.
- ② Learn the classifier h_2 on subset $S_2 \subset S \setminus S_1$ with half of elements of S_2 wrongly classified by h_1 .
- ③ Learn classifier h_3 on subset $S_3 \subset S \setminus S_1 \setminus S_2$: contains elements for which rules h_1 and h_2 answer differently.
- ④ H : Majority vote between answers of h_1 , h_2 and h_3 .

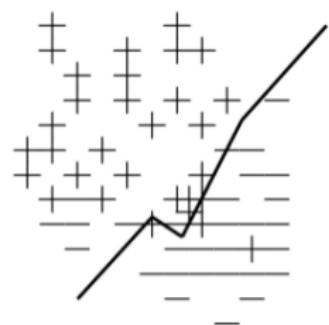
A toy example (3)



h_2 learned on S_2



S_3 and h_3



S and H

1 Boosting algorithm

- Idea
- A first simple example to understand boosting idea
- **Probabilistic boosting and Adaboost**
- A simple example for understanding Adaboost

2 Application to face detection

3 main ideas :

- ① A set of specialized experts and ask them to vote to take a decision.
- ② Adaptive weighting of votes by multiplicative update.
- ③ Modifying example distribution to train each expert, increasing the weights iteratively of examples misclassified at previous iteration.

- S learning dataset
- Initialisation : all samples have same weights ($D_1(i), i \in \{1 \dots m\}$).
- Iterations : for $t \in \{1 \dots T\}$
 - learn h_t by minimisation of an **error**
 - compute **weight** α_t and **weights** $D_t(i)$ (or distribution)
- Compute strong classifier :

$$\text{sign} \left(H(x) = \sum_{t=1}^T \alpha_t h_t(x) \right)$$

Adaboost : Adapative boosting

- $S = \{(x_1, y_1), \dots, (x_m, y_m)\}$ with $x_i \in X$ and $y_i \in \{-1, +1\}$
- Initialisation : $D_1(i) = \frac{1}{m}$ with $i \in \{1 \dots m\}$
- For $t \in \{1 \dots T\}$:
 - **find** $h_t : X \rightarrow \{-1, +1\}$ minimizing error ϵ_t defined by :

$$\epsilon_t = \sum_{i=1}^m D_t(i)[y_i \neq h_t(x_i)]$$

- **compute weights** :

$$\alpha_t = \frac{1}{2} \ln\left(\frac{1 - \epsilon_t}{\epsilon_t}\right)$$

- **compute distribution** :

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(i))}{Z_t}$$

where Z_t is for normalisation : $1 = \sum_{i=1}^m D_t(i)$

- **Compute strong classifier** :

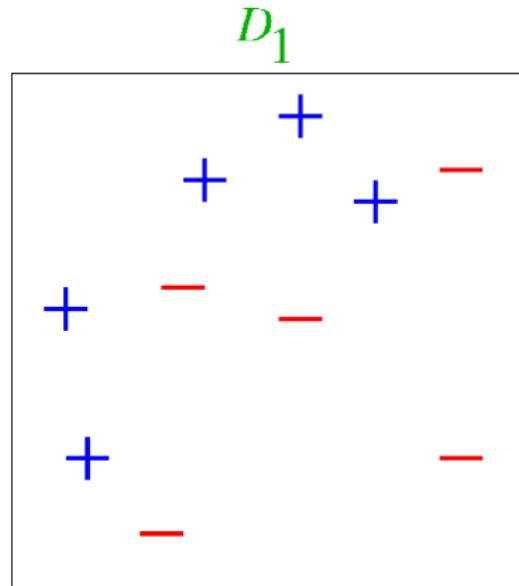
$$\text{sign} \left(H(x) = \sum_{t=1}^T \alpha_t h_t(x) \right)$$

1 Boosting algorithm

- Idea
- A first simple example to understand boosting idea
- Probabilistic boosting and Adaboost
- A simple example for understanding Adaboost

2 Application to face detection

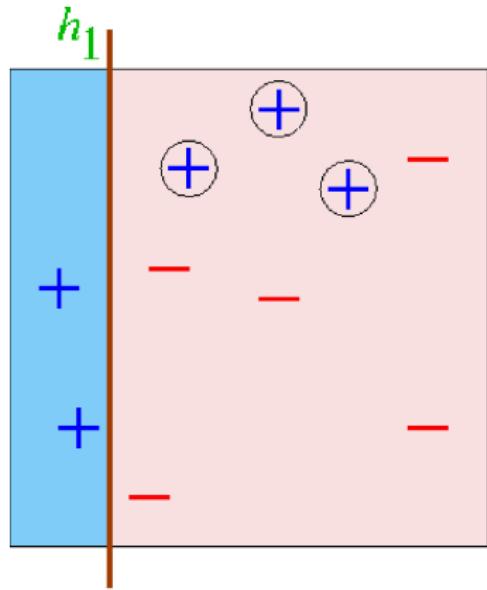
Initialisation



$D_1 :$

0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----

Step 1



$$\epsilon_1 = 0.30$$

$$\alpha_1 = \frac{1}{2} \ln\left(\frac{1 - \epsilon_1}{\epsilon_1}\right) = \frac{1}{2} \ln\left(\frac{0.7}{0.3}\right) = 0.42$$

$$D_2(1) = \frac{1}{Z_1} D_1(1) e^{\alpha_1} = \frac{0.152}{Z_1}$$

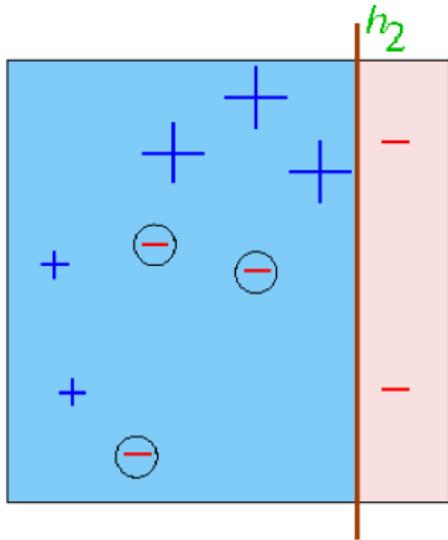
$$D_2(4) = \frac{1}{Z_1} D_1(4) e^{-\alpha_1} = \frac{0.065}{Z_1}$$

$$Z_1 = 3 \times 0.152 + 7 \times 0.065 = 0.911$$

$D_2 :$

0.167	0.167	0.167	0.071	0.071	0.071	0.071	0.071	0.071	0.071
-------	-------	-------	-------	-------	-------	-------	-------	-------	-------

Step 2



$$\epsilon_2 = 0.21$$

$$\alpha_2 = \frac{1}{2} \ln\left(\frac{1 - \epsilon_2}{\epsilon_2}\right) = 0.65$$

$$D_3(1) = \frac{D_2(1)}{Z_2} e^{-\alpha_2} = 0.167 e^{-0.65} = \frac{0.0876}{Z_2}$$

$$D_3(4) = \frac{D_2(4)}{Z_2} e^{-\alpha_2} = 0.071 e^{-0.65} = \frac{0.036}{Z_2}$$

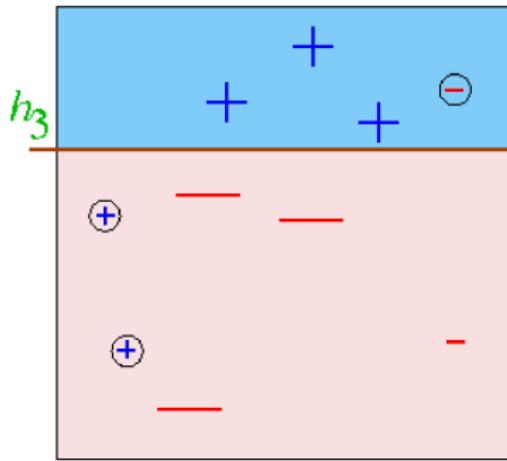
$$D_3(8) = \frac{D_2(8)}{Z_2} e^{+\alpha_2} = 0.071 e^{+0.65} = \frac{0.1357}{Z_2}$$

$$Z_2 = 3 \times 0.0876 + 4 \times 0.036 + 3 \times 0.1357 = 0.814$$

$D_3 :$

0.11	0.11	0.11	0.044	0.044	0.044	0.044	0.17	0.17	0.17
------	------	------	-------	-------	-------	-------	------	------	------

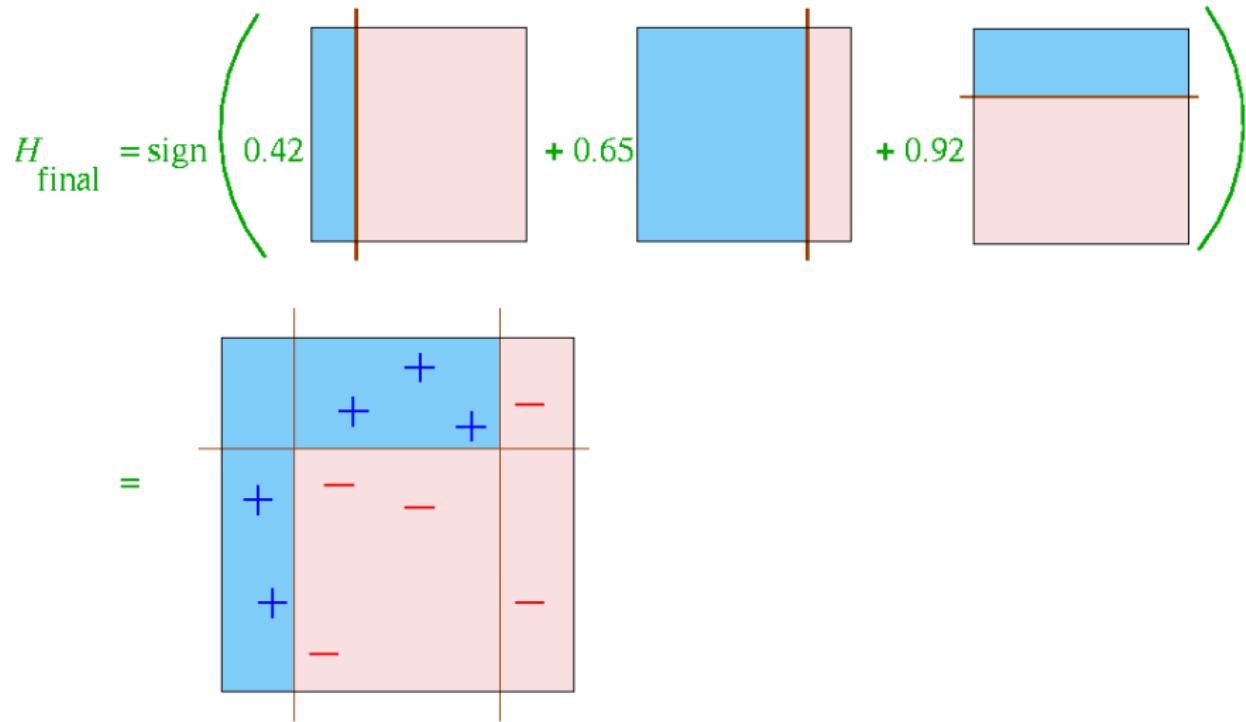
Step 3



$$\epsilon_3 = 3 \times 0.044 = 0.13$$

$$\alpha_3 = 0.92$$

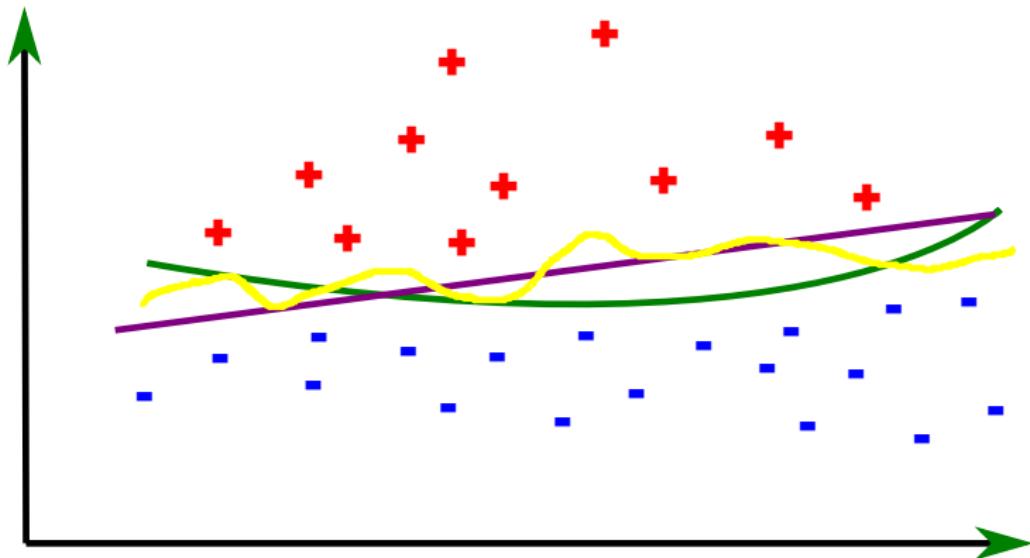
Final classifier



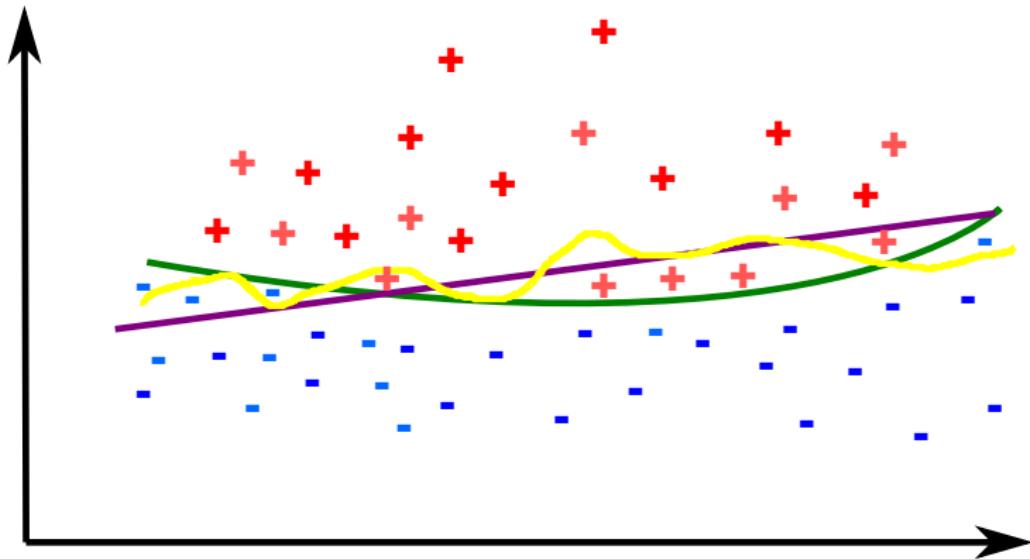
Advantages and drawbacks

- Advantages :
 - Fast
 - Easy to implement
 - only 1 parameter : number of boosting iterations
 - extensible to multi classes classification
 - able to detect outliers
- Drawbacks :
 - depends on learning data and weak classifiers
 - could fail if :
 - the weak classifier is too complex
 - low margins → overfitting
 - the weak classifier is too weak
 - underfitting
 - noise sensitive

Overfitting and Underfitting



Overfitting and Underfitting



- Solve overfitting when noisy data :
 - Weight Decay (1998) : introduces slack variables regularizing the error
 - GentleBoost (1998)
 - LogitBoost (2000) : additive logistic regression model (statistical point of view)
 - Regularized AdaBoost (2000) : “soft” margins
 - BrownBoost (2001) : removes examples many times misclassified, considered as “noisy”
 - WeightBoost (2003) : weights in final decision dependent on input
- Reduce the number of features (or Weak Learner) :
 - FloatBoost (2003) : removes the worst WL after each iteration
 - JointBoost (2004) : multi-class, jointly train N binary classifier sharing same features
- Adaboost Multi-class
 - AdaBoost.M1 : same as standard but $y \in \{1 \dots k\}$
 - AdaBoost.M2 : introduces a confidence degree
 - AdaBoost SAMME (2006)

Multi-class boosting : AdaBoost SAMME

- $S = \{(x_1, y_1), \dots, (x_m, y_m)\}$ with $x_i \in X$ and $y_i \in \{1 \dots k\}$
- Initialisation : $D_1(i) = \frac{1}{m}$ with $i \in \{1 \dots m\}$
- For $t \in \{1 \dots T\}$:
 - **find** $h_t : X \rightarrow \{1 \dots k\}$ minimizing error ϵ_t defined by :

$$\epsilon_t = \sum_{i=1}^m D_t(i) g(y_i, h_t(x_i)) \text{ with } g(a, b) = 1 \text{ if } a = b \text{ else } 0$$

- **compute weights** :

$$\alpha_t = \frac{1}{2} \ln\left(\frac{1 - \epsilon_t}{\epsilon_t}\right) + \frac{1}{2} \ln(k - 1)$$

- **compute distribution** :

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t g(h_t(i), y_i))}{Z_t} \text{ where } 1 = \sum_{i=1}^m D_t(i)$$

- Compute **strong classifier** :

$$\arg \max_k \left(\sum_{t=1}^T \alpha_t g(h_t(x), k) \right)$$

- Detection, tracking and face recognition
- Spam, Zip Code OCR
- Text classification : Schapire and Singer
- OCR : Schwenk and Bengio (neural networks)
- Natural language Processing : Collins ; Haruno, Shirai and Ooyama
- Image retrieval : Thieu and Viola
- Medical diagnosis : Merle et al.
- Fraud Detection : Rätsch & Müller 2001
- Drug Discovery : Rätsch, Demiriz, Bennett 2002
- Elect. Power Monitoring : Onoda, Rätsch & Müller 2000

Example using scikit-learn

From <https://scikit-learn.org/stable/modules/ensemble.html#gradient-tree-boosting> :

```
from sklearn.model_selection import cross_val_score
from sklearn.datasets import load_iris
from sklearn.ensemble import AdaBoostClassifier

iris = load_iris()
clf = AdaBoostClassifier(n_estimators=100)
scores = cross_val_score(clf, iris.data, iris.target, cv=5)
scores.mean()
```

1 Boosting algorithm

- Idea
- A first simple example to understand boosting idea
- Probabilistic boosting and Adaboost
- A simple example for understanding Adaboost

2 Application to face detection

Face detection



Face detection



How to ?



A composite image showing a woman speaking into a microphone on the left and a presentation slide on the right.

The presentation slide has a dark header bar with the text "A deeper view of Deep Learning". Below the header, the names "Mélanie Ducoffe, Frédéric Precioso" are listed. Underneath that, the text "Laboratoire i3S - UNS CNRS UMR 7271" is displayed. At the bottom of the slide, the date "06/07/2016" is shown. The footer of the slide features logos for CNRS, i3S, and Université Paul-Sabatier Toulouse.

How to ?



The image is a composite of two parts. On the left, a woman with glasses and dark hair is speaking into a black microphone. On the right, there is a presentation slide with a dark blue header containing the text "A deeper view of Deep Learning". Below the header, the names "Mélanie Ducoffe, Frédéric Precioso" are listed. Further down, the text "Laboratoire I3S - UNS CNRS UMR 7271" is displayed, along with the date "06/07/2016". At the bottom of the slide, there are logos for CNRS (a purple circle with white text), I3S (a purple stylized logo with yellow lines), and Université Grenoble Alpes (a green and yellow logo).

A red square box is drawn around the top-left corner of the woman's image.

M

E

CNRS

i3S

Université
Grenoble
Alpes

A deeper view of Deep Learning

Mélanie Ducoffe, Frédéric Precioso

Laboratoire I3S - UNS CNRS UMR 7271

06/07/2016

How to ?



The image is a composite of two parts. On the left, a woman with glasses and brown hair is speaking into a black microphone. A red square box is drawn around her head area. On the right, there is a presentation slide with a dark blue header containing the white text "A deeper view of Deep Learning". Below the header, the names "Mélanie Ducoffe, Frédéric Precioso" are listed. Further down, the text "Laboratoire I3S - UNS CNRS UMR 7271" is displayed, followed by the date "06/07/2016". At the bottom of the slide, there are logos for CNRS (a purple circle with "CNRS" in white), I3S (a purple stylized "i3S" logo with yellow lines), and Université Grenoble Alpes (a logo with "Université" in blue, "Grenoble" in green, and "Alpes" in orange).



A composite image showing a woman speaking into a microphone on the left and a presentation slide on the right. A red box highlights a portion of the woman's face in the video feed.

A deeper view of Deep Learning

Mélanie Ducoffe, Frédéric Precioso

Laboratoire i3S - UNS CNRS UMR 7271

06/07/2016

M E

CNRS i3S Université Paul-Sabatier Toulouse

How to ?



M

A deeper view of Deep Learning

Mélanie Ducoffe, Frédéric Precioso

Laboratoire I3S - UNS CNRS UMR 7271

06/07/2016

E

CNRS

i3S

Université
Nice
Savoie
Toulon

Problem description

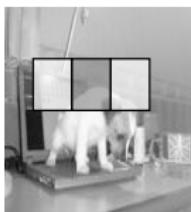
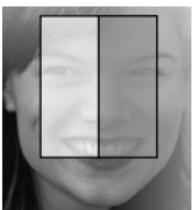
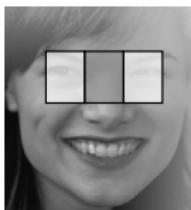
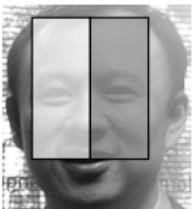
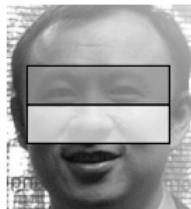
- slide a window across image and evaluate a face model at every location.
- sliding window detector must evaluate tens of thousands of location/scale combinations.
- faces are really rare compare to the number of evaluations
 - for computational efficiency, we should try to spend as little time as possible on the non-face windows
 - false positive rate should be very low ($< 10^{-6}$ for an image of 1 Mpixels)

- real-time object detection
- slow training but fast detection
- 3 main ideas :
 - integral images for fast feature evaluation
 - boosting for feature selection
 - cascade classifiers for fast rejection of non face

[CVPR 2001] : P. Viola and M. Jones. "Rapid object detection using a boosted cascade of simple features".

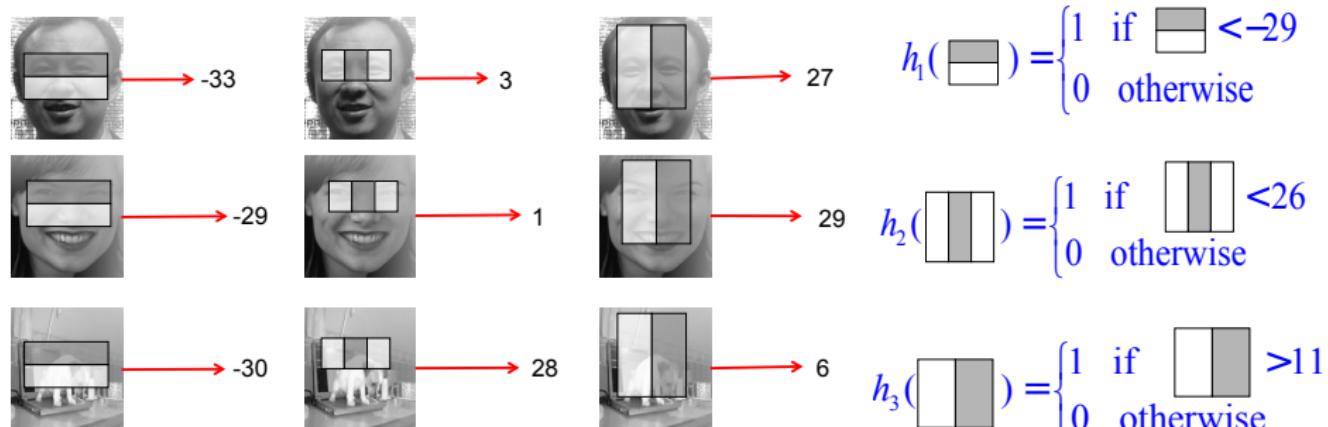
[IJCV 2004] : P. Viola and M. Jones. "Robust real-time face detection", 57(2).

Finding weak classifiers for faces



$$f_t(W) = \sum_{\text{white}} \text{pixels} - \sum_{\text{black}} \text{pixels}$$

How to learn classifiers ?



Generalisation :

$$h_t(W) = \begin{cases} 1 & \text{if } p_t f_t(W) > p_t \theta_t \\ 0 & \text{else} \end{cases}$$

where the parity p_t equals ± 1 , threshold θ_t has to be learned.

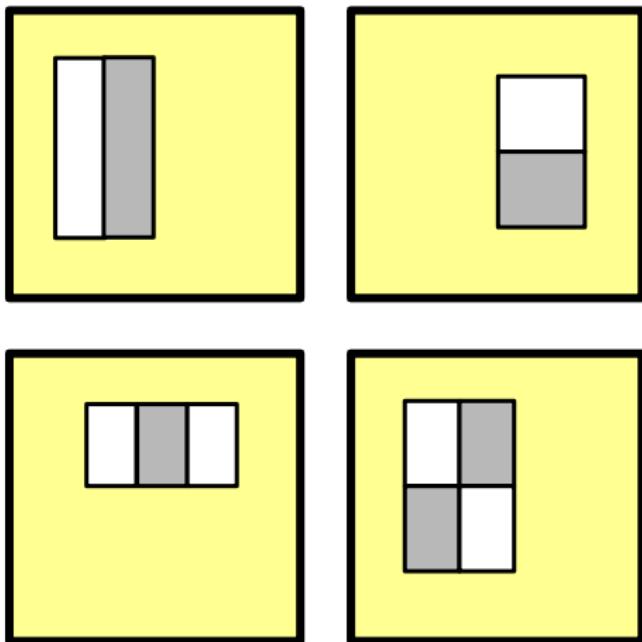
- S learning dataset : faces and non faces
- Initialisation : same weights for faces ($\frac{1}{m_v}$) and same weights for non faces ($\frac{1}{m_{nv}}$).
- Iterations : for $t \in \{1 \dots T\}$
 - for all available filters :
 - learn threshold θ for all learning samples
 - select best filter h_t with threshold θ_t
 - compute weights α_t and $D_t(i)$
 - compute strong classifier

K : number of filters

Learning complexity : $O(T \cdot (\text{nb. samples}) \cdot K)$

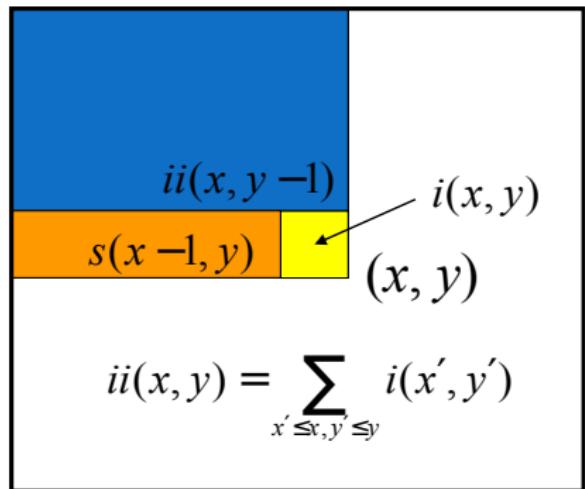
Weak classifiers : rectangular filters

- based on Haar filters
- difference between sum of white related pixels and sum of black related pixels
- 3 types : 2, 3 or 4 rectangles
- combinaisons (rectangles, sizes, positions) for a 24x24 window : 160 000 filters



Integral Image (II) computation

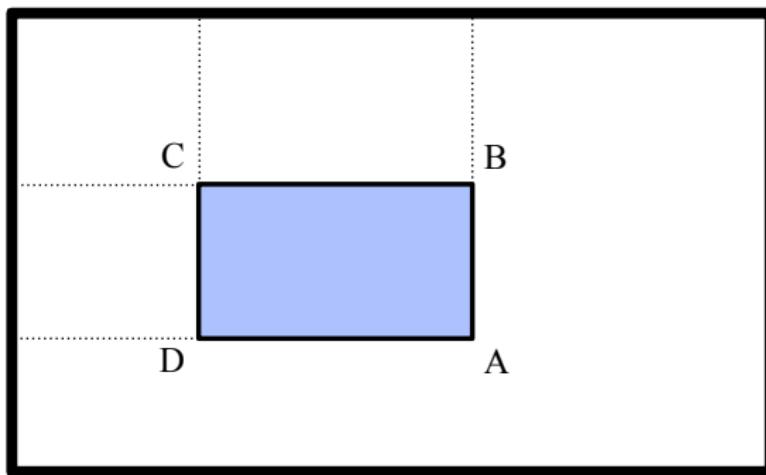
- i is the initial image; ii is the integral image
- Each pixel (x,y) of the integral image represents the sum of all pixels above and to the left of (x,y) , from the initial image.
- Integral image is computed in one pass using a temporary structure representing the sum of pixels on the same row, at left : s .



$$\begin{aligned}s(x, y) &= s(x - 1, y) + i(x, y) \\ ii(x, y) &= ii(x, y - 1) + s(x, y)\end{aligned}$$

Use of integral image

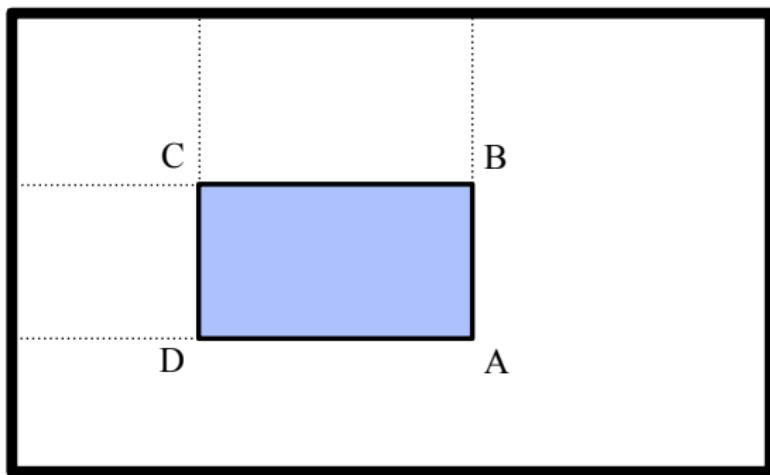
Compute the sum of pixels inside a rectangular area ABCD :



$$S = \sum_{x=x_C}^{x_A} \sum_{y=y_D}^{y_C} i(x, y) = ii_A - ii_B + ii_C - ii_D$$

Use of integral image

Compute the sum of pixels inside a rectangular area ABCD :



$$S = \sum_{x=x_C}^{x_A} \sum_{y=y_D}^{y_C} i(x, y) = ii_A - ii_B + ii_C - ii_D$$

Only 3 operations are needed !

Other optimisations :

- Different scales :
 - Instead of changing samples scales, we change the filters scales.
- Normalisation of images computed from integral images and squared integral images :

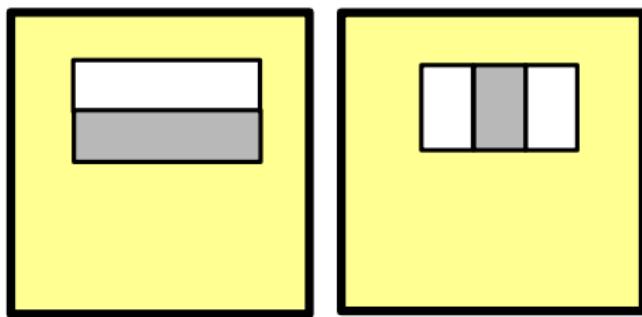
$$\sigma^2 = m^2 - \frac{1}{N} \sum x^2$$

Learning dataset :

- 5000 faces
 - front oriented
 - size 24x24 pixels
 - normalised
- 300 millions of non faces
 - taken from 9500 images without faces
- variability :
 - between individuals
 - illumination
 - pose

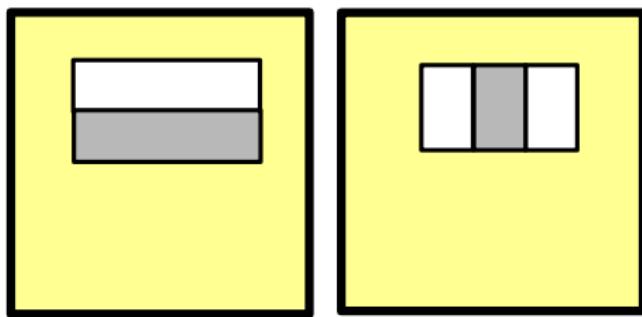


Selected filters by *boosting* : the first 2 features



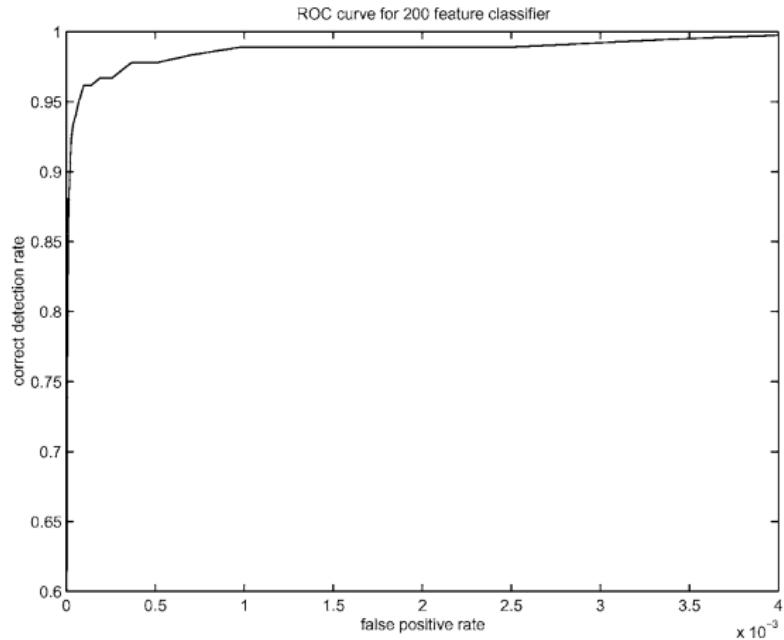
- conduct to :
 - 100% detection rate (ratio true positives over positives)
 - 50% false positives rate (ratio false positives over negatives)
- able to eliminate non faces very fastly

Selected filters by *boosting* : the first 2 features



- conduct to :
 - 100% detection rate (ratio true positives over positives)
 - 50% false positives rate (ratio false positives over negatives)
 - able to eliminate non faces very fastly
- ⇒ Fast elimination of non faces and increase of filters complexity

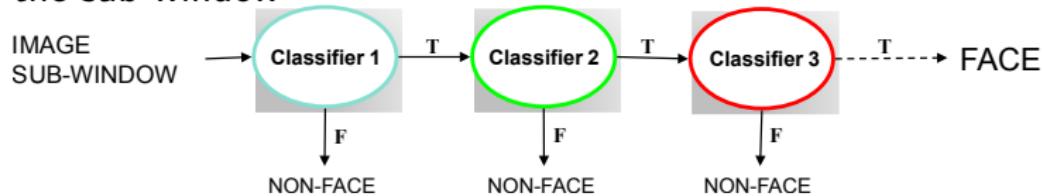
The 200 best filters selected : ROC curve



A 200-feature classifier can yield 95% detection rate and a false positive rate of 1 in 14084. To be practical for real application, the false positive rate must be closer to 1 in 1,000,000.

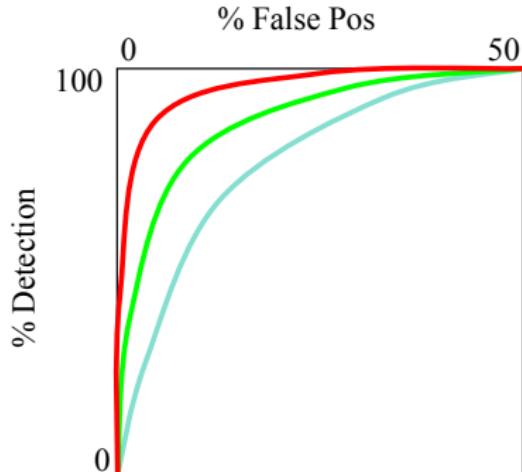
Attentional Cascade

- We start with simple classifiers which reject many of the negative sub-windows while detecting almost all positive sub-windows
- Positive response from the first classifier triggers the evaluation of a second (more complex) classifier, and so on
- A negative outcome at any point leads to the immediate rejection of the sub-window

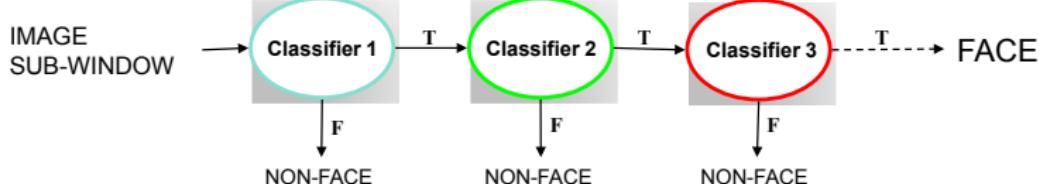


Attentional Cascade

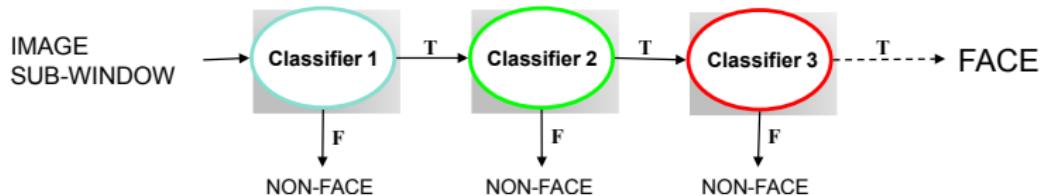
ROC Curve



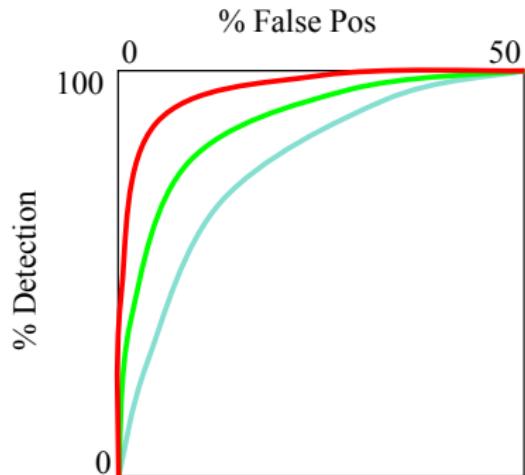
Chain classifiers that are progressively more complex and have lower false positive rates



Cascade structure



ROC Curve



$$F = \prod_{i=1}^K f_i$$

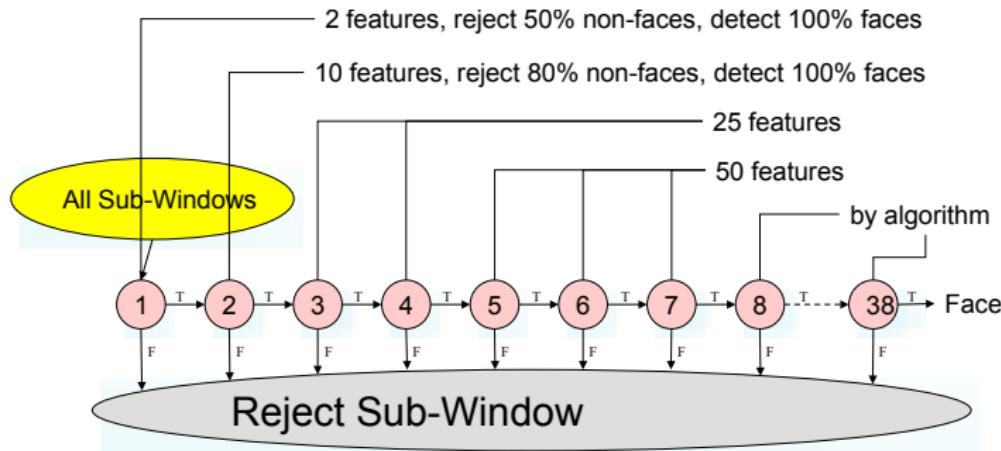
$$D = \prod_{i=1}^K d_i$$

A detection rate of 0.9 and a false positive rate on the order of 10^{-6} can be achieved by a 10-stage cascade if each stage has a detection rate of 0.99 ($0.99^{10} \approx 0.9$) and a false positive rate of about 0.30 ($0.3^{10} \approx 6 \cdot 10^{-6}$)

Algorithm for determining the stages of the cascade

- choose $f = \max_i f_i$ and $d = \min_i d_i$
- choose F
- P : set of faces
- N : set of non faces
- initialisations : $F_0 = 1.0 \quad D_0 = 1.0 \quad i = 0$
- while $F_i > F$
 - add a new stage : $i \leftarrow i + 1$
 - $n_i = 0 \quad F_i = F_{i-1}$
 - while $F_i > f F_{i-1}$:
 - add filters : $n_i \leftarrow n_i + 1$
 - boosting with n_i filters
 - evaluation of cascade classifier on validation data : F_i and D_i
 - lower the threshold of classifier i for having a detection rate at least : $d.D_{i-1}$ (impacts also F_i)
 - empty N
 - false positives go to N for the next step

The final cascade

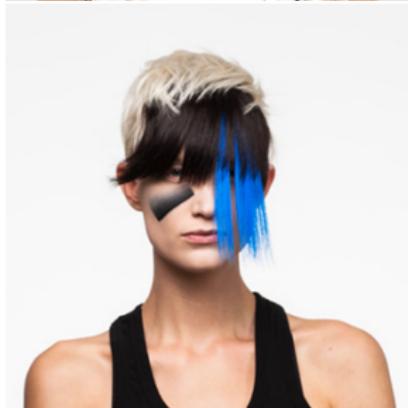
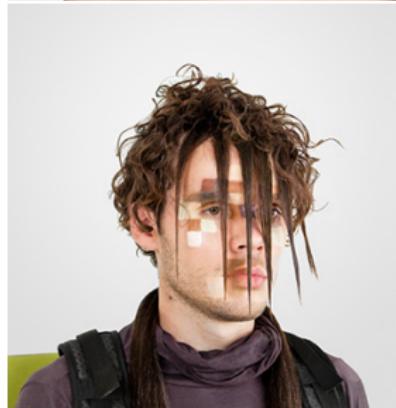
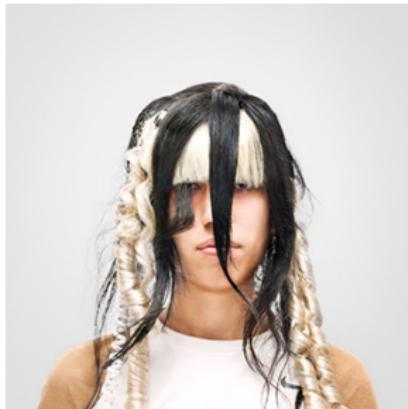


- Consists in combining in cascade more and more complex classifiers
 - 38 stages
 - total of 6060 classifiers
- Historical performances :
 - on a Pentium III, 700 MHz, processing of a 384x288 pixels images in 0.067 seconds (15 Hz)

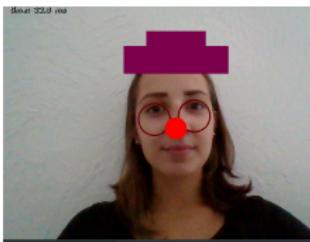
- Other filters (rotation 45 deg)
- Nested cascades :
 - eyes detection
 - nose detection
 - smile detection

How to escape to facedetect ?

see <https://cvdazzle.com/>



Let's play with facedetect !



Alexia



Erwan



Anthony



Alexis G.



Yilei



Simon



Vincent



Maxime



Mariana

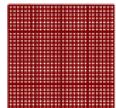
Tests on MNIST database

- 28x28 grey level images : 784 pixels
- Haar filters : 'type-2-x' and 'type-2-y' : 158760 filters
- if we remove filters of diagonal size $< \sqrt{50}$ and take only 1 over 128 filters : 859 filters
- only two classes : 4 and 8 ; confusion matrix $\begin{bmatrix} 980 & 2 \\ 6 & 968 \end{bmatrix}$
- with all classes

887	0	15	7	7	7	11	1	38	7
1	1105	8	3	0	5	1	1	11	0
6	10	784	52	44	9	67	5	52	3
2	2	76	805	2	26	2	11	48	36
5	1	13	7	829	5	12	7	10	93
3	2	16	104	1	666	23	13	42	22
12	4	27	0	31	13	867	0	4	0
1	9	24	12	9	7	0	820	17	129
56	12	21	86	10	28	18	8	680	55
4	8	3	20	120	4	0	138	24	688

Tests on MNIST database

- only two classes : 4 and 8



- haar features with resolution quarter w(=28) and quarter h
- size of descriptor : 168 (data 784 pixels)
- confusion matrix

$$\begin{bmatrix} 971 & 11 \\ 12 & 962 \end{bmatrix}$$

- with all classes

- confusion matrix

825	0	18	4	10	49	40	2	31	1
0	1042	6	2	0	1	3	17	64	0
15	45	812	11	22	4	20	17	80	6
52	60	51	653	4	46	1	27	82	34
5	4	1	3	838	5	15	11	21	79
42	19	17	116	27	467	15	8	123	58
46	4	45	0	20	22	757	5	58	1
1	22	22	12	30	21	0	767	22	131
27	19	29	58	28	26	15	8	744	20
5	15	5	22	439	8	2	40	37	436

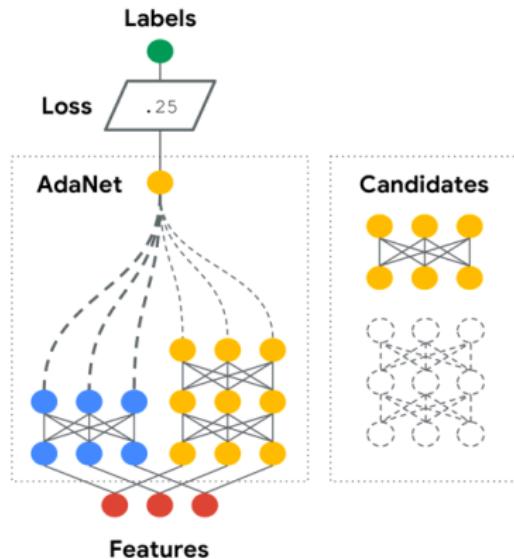
Pixels as image feature

- Confusion matrix

883	0	25	3	4	28	23	3	3	8
0	1070	3	8	3	1	4	25	21	0
30	35	596	32	18	8	208	25	75	5
28	32	19	678	2	92	30	33	74	22
4	2	17	14	708	16	10	80	35	96
29	32	6	122	27	526	22	19	71	38
20	10	35	6	26	32	822	1	6	0
7	16	23	8	14	7	1	804	20	128
40	48	11	91	15	34	22	18	661	34
9	11	23	32	161	18	1	169	34	551

- does it means sens ?
- is it better ? worse ?

boosting and deep learning



AdaNet :

<https://github.com/tensorflow/adanet>